

# Revisiting Hall-Petch strengthening in FCC multi-principal element alloys: what grain size, composition, and processing can and cannot explain

Mrinalini Mulukutla<sup>a,\*</sup>, Shakti Prasad Padhy<sup>a,b</sup>, Wenle Xu<sup>a</sup>, Bibhu Prasad Sahu<sup>a</sup>, Joydeep Kundu<sup>a</sup>, Vasanth C. Shunmugasamy<sup>a</sup>, Douglas Allaire<sup>b</sup>, Ibrahim Karaman<sup>a</sup>, Raymundo Arróyave<sup>a,b</sup>

<sup>a</sup>Department of Materials Science and Engineering, Texas A&M University, College Station, TX, 77843, USA

<sup>b</sup>J. Mike Walker '66 Department of Mechanical Engineering, Texas A&M University, College Station, TX, 77843, USA

---

## Abstract

Grain size is one of the primary microstructural variables controlling strength, but its role in FCC multi-principal element alloys cannot be separated easily from the effects of chemistry and processing. In this work, we examine the relationships among grain size, composition, processing, yield strength, and Vickers hardness using 94 alloy candidates representing 82 compositions and six experimental batches, spanning grain sizes of 15–212  $\mu\text{m}$  and yield strengths of 152–544 MPa. The classical Hall–Petch relation provides a reproducible yield-strength baseline, giving  $R^2 = 0.41$  under record-level and 0.37 under batch-level validation, while the inclusion of composition raises these values to 0.65 and 0.63 with a shared slope of 766 MPa· $\mu\text{m}^{1/2}$ . The strongest low-complexity model also contains a coupled contribution from mean grain size and its within-map standard deviation, reaching 0.69 under both protocols. However, the present data cannot separate the physical origin of this contribution from the covariance among composition, processing, and microstructure. The tested solid-solution-strengthening estimates and nonlinear regressors do not improve batch-level predictions once the measured inputs are included. Hardness shows a different dependence and does not provide a transferable substitute for yield strength across the campaigns. The comparison with a heterogeneous literature dataset shows that transfer across studies will require prospective datasets in which composition, processing history, microstructure, test conditions, and other experimental factors are documented and controlled consistently.

**Keywords:** High-entropy alloys, Yield strength, Vickers hardness, Solid solution strengthening, Symbolic regression, Cross-validation

---

## Abbreviations

**AI** Artificial intelligence  
**AIC** Akaike information criterion  
**AICc** Corrected Akaike information criterion  
**ARMOTE-CV** Automated Regression with Multi-Objective Tuning and Evaluation via Cross-Validation  
**AUC** Area under the receiver-operating-characteristic curve  
**BIC** Bayesian information criterion  
**BIRDSHOT** Batch-wise Improvement in Reduced Materials Design Space Using a Holistic Optimization Technique  
**BO** Bayesian optimization  
**CV** Cross-validation  
**EBS** Electron backscatter diffraction  
**EDS** Energy-dispersive X-ray spectroscopy  
**EML** Binary symbolic-regression operator,  $\text{eml}(x, y) = \exp(x) - \ln(y)$   
**FCC** Face-centered cubic  
**HP** Hall-Petch  
**HTMDEC** High-Throughput Materials Discovery for Extreme Conditions  
**HV** Vickers hardness  
**JMAK** Johnson-Mehl-Avrami-Kolmogorov  
**KRR** Kernel ridge regression

**Lasso** Least absolute shrinkage and selection operator  
**LightGBM** Light Gradient Boosting Machine  
**LOBO** Leave-one-batch-out  
**LOO** Leave-one-out  
**ML** Machine learning  
**MLP** Multilayer perceptron  
**MPEA** Multi-principal element alloy  
**OLS** Ordinary least squares  
**PCA** Principal component analysis  
**PSIS-LOO** Pareto-smoothed importance-sampling leave-one-out  
**PySR** Python symbolic-regression package  
**RBF** Radial basis function  
**RMSE** Root-mean-square error  
**SD** Standard deviation  
**SHAP** SHapley Additive exPlanations  
**SI** Supplementary Information  
**SISSO** Sure independence screening and sparsifying operator  
**SSS** Solid-solution strengthening  
**SVR** Support vector regression  
**VEC** Valence electron concentration  
**VLC** Varvenne-Leyson-Curtin  
**XGBoost** Extreme Gradient Boosting  
**YS** Yield strength

---

\*Corresponding author.

Email address: [mrinalini.mulukutla@tamu.edu](mailto:mrinalini.mulukutla@tamu.edu)  
(Mrinalini Mulukutla)

## 1. Introduction

Grain size is one of the primary microstructural variables controlling strength in FCC alloys, and its contribution is usually described by the Hall-Petch relation

$$\sigma_y = \sigma_0 + k_{\text{HP}}d^{-1/2} \quad (1)$$

which separates grain-boundary strengthening from an intercept  $\sigma_0$  [1, 2, 3]. The intercept is the strength that remains in the large-grain limit, where lattice friction and solute strengthening set the resistance to slip. In concentrated FCC alloys, composition acts on both terms, changing solute resistance, ordering, and stacking fault energy, and altering boundary-mediated deformation. Both therefore vary among alloy systems [4, 5, 6]; reported slopes span 265 to 1244 MPa· $\mu\text{m}^{1/2}$  among concentrated FCC alloys [5, 7]. Several other grain-size laws fit a limited range of  $d$  almost as well as Eq. (1), so a good regression fit identifies the trend while leaving the mechanism open.

Physics-based solid solution strengthening models estimate the solute contribution to the intercept from atomic size and modulus misfit, while composition descriptors summarize composition through quantities such as valence electron concentration and atomic size mismatch [8, 9, 10, 11, 12]. Both provide compact alternatives to using unconstrained elemental fractions. However, these descriptors are calculated directly from composition and may add little information after composition is already included. Non-linear and symbolic models can capture interactions that a linear model misses, but they also search a much larger space. As a result, they can exploit campaign identity, correlated descriptors, or unstable denominators. A feature can be physically meaningful without being transferable, and an equation can be interpretable without being a validated physical law.

Hardness serves as an additional, complementary target. Vickers hardness captures the material’s resistance to deformation under localized loading. The empirical Tabor relation relates hardness to a representative flow stress, rather than directly to 0.2% offset YS [13, 14]. Applying the same modeling and validation framework to YS and HV shows whether the relevant features depend on the property being modeled.

The alloys examined here were selected by closed-loop design campaigns to optimize target properties. Multi-principal element alloys span composition and processing spaces that are high dimensional and too large to explore exhaustively [15, 16, 17, 18]. Closed-loop and interpretable machine learning (ML) frameworks can speed up this search [19, 20, 21, 22, 23, 24]. However, a property model becomes harmful when a good interpolation score is used to support decisions about new compositions, processing routes, or experimental campaigns. This risk is especially high for small and sparse experimental datasets that represent these very high dimensional design spaces. Nearby compositions, repeated conditions, and campaign-specific processing can make a random train/test split much easier than the intended alloy design problem [25, 26, 27].

Hall-Petch parameters for FCC alloys are usually reported a few compositions at a time. Of the 630 alloys compiled by Borg et al., 27 single-phase FCC compositions carry grain size

together with yield strength, drawn from 25 separate studies [28]. Six carry grain size, yield strength and hardness together. The 81 compositions measured here carry all three, under one processing route and one characterization protocol. This is the largest such set reported for single-phase FCC multi-principal element alloys, and it lets composition, processing and grain size be separated without correcting across studies.

Figure 1 and Table 1 define the five model families used throughout this paper, with the same names and order: Family 1 (Classical Hall-Petch), Family 2 (Physics-derived descriptors), Family 3 (Composition/processing), Family 4 (Non-linear estimators), and Family 5 (Symbolic regression). The families ask five sequential questions: What does grain size explain? Do physics-based descriptors capture the remaining strength? What do measured composition and processing add? Does non-linearity help when inputs are matched? Can symbolic search produce useful closed forms? A higher family number does not imply greater physical fidelity. We combine this hierarchy with 5-fold CV, record-level LOO, batch-level LOBO, an external literature stress test, and singularity auditing of closed-form equations.

Table 1: Five model families used in the main paper and Supplementary Information.

| Family | Fixed name                  | Analyses                                                                             |
|--------|-----------------------------|--------------------------------------------------------------------------------------|
| 1      | Classical Hall-Petch        | Classical and alternative grain-size laws                                            |
| 2      | Physics-derived descriptors | SSS models, Wen descriptors, and PCA-OLS                                             |
| 3      | Composition/processing      | OLS and composition-processing Hall-Petch models, including grain-width interactions |
| 4      | Non-linear estimators       | Matched linear and non-linear estimators                                             |
| 5      | Symbolic regression         | PySR, SISSO, and fixed closed forms                                                  |

## 2. Dataset and descriptors

### 2.1. Alloy Synthesis and characterization

The dataset contains 94 FCC alloy conditions in the Al–Co–Cr–Cu–Fe–Mn–Ni–V system. YS is available for 93 conditions, and HV is available for all 94. Samples carry a three-letter batch code. The first two letters identify the design campaign and the third the batch within it, so AA, BB and CB denote the first, second and third BIRDSHOT campaigns. An earlier Campaign AA established the closed-loop, multi-objective alloy-discovery workflow used across this effort [20]; its records are not part of the present property modeling dataset because that campaign used a different processing route. The samples analyzed here were generated from the subsequent BB and CB Bayesian optimization campaigns at Texas A&M University within BIRDSHOT Center, supported by the U.S. Army DEVCOM High-Throughput Materials Discovery for Extreme Conditions (HT-MDEC) initiative [21, 22]. The BB campaign contributes 35 candidates from batches BBA, BBB, and BBC, including 34 with YS. It explored composition under an FCC phase stability constraint while holding processing fixed. The complete




Figure 1: Five family framework used throughout the paper: Classical Hall-Petch, Physics-derived descriptors, Composition/processing, Non-linear estimators, and Symbolic regression. S1–S4 name the matched input sets in Family 4; individual inputs, including  $SD_{\text{grain}}$ , are also tested in Families 2, 3, and 5. Internal evaluation uses 5-fold CV, LOO, and LOBO; external literature testing and singularity auditing provide additional challenge and equation safety checks.

BB-campaign synthesis and specimen preparation workflow is documented in Ref. [22].

The CB campaign contributes 59 candidates from CBA, CBB, and CBC. Building on the BB-campaign results, it focused on a V–Cr–Mn–Fe–Co–Ni design space and jointly varied composition with three processing variables: cold work, recrystallization temperature, and hold time. Candidate alloys were iteratively selected from an FCC-constrained, manufacturable design space with the models updated after each iteration using the newly acquired data. A modified JMAK-type static-recrystallization and grain-growth prior supplied grain size estimates over the feasible design space to guide candidate selection. That model and the complete CB-campaign workflow will be reported in a separate future publication. The property models evaluated here use grain sizes measured by electron backscatter diffraction (EBSD), not these design-stage predictions.

Overall, both campaigns used the same high level production route: high-purity elemental feedstock was vacuum arc melted, homogenized, cold rolled, recrystallized, and quenched. Energy-dispersive X-ray spectroscopy (EDS) and X-ray diffraction were used to check realized chemistry and phase constitution, while electron microscopy and EBSD documented the microstructure. Grain size was measured from EBSD maps on surfaces normal to the rolling direction and summarized by the equivalent circle mean  $d$  and within-map standard deviation  $SD_{\text{grain}}$ . Representative candidate alloy verification is shown in Figure 2. Quasistatic tensile testing was used to determine 0.2% offset YS, while Vickers indentation tests provided the HV. Further details for BBA-BBC are given in Ref. [22]; a detailed account of CBA-CBC will follow in the campaign publication.

## 2.2. Coverage, dependence, and feature sets

The 94 candidates in this study represent 82 unique compositions: 73 chemistries occur once, and nine repeat across 21

records. Processing differs within eight of those nine groups, so these repeated chemistries sample grain size at fixed chemistry under varying processing conditions rather than serving as replicates; the exception is the CBB04/CBB12 pair, which shares both composition and processing (60% cold work, 875 °C, 1 h) and differs by 32 MPa in YS and 4.7 in HV, the only direct estimate of repeatability the dataset contains. The 93-record YS subset contains 81 unique compositions. One composition appears in both CBB and CBC, so LOBO reduces campaign dependence without removing all exact-composition overlap, and LOO retains an identical chemistry in the training set for 21 of the 94 records. Mean  $d$  ranges from 14.66 to 211.75  $\mu\text{m}$ ,  $SD_{\text{grain}}$  from 8.32 to 321.06  $\mu\text{m}$ , YS from 151.5 to 544.5 MPa, and HV from 57.6 to 227.5. Because  $d$  and  $SD_{\text{grain}}$  are strongly correlated ( $r = 0.801$ ), the standard deviation carries much of the same information as the mean. We report it as a measure of distribution width. Measurement scatter is a separate quantity:  $SD_{\text{YS}}$  is recorded for 80 conditions (0.2–98.7 MPa) and  $SD_{\text{HV}}$  for all 94 (1.43–12.44 HV). These measurement standard deviations are not model inputs because replicate counts and reporting definitions are not consistent across the dataset. The Bayesian optimization campaigns sample design-relevant regions, not a random composition simplex (Figure 3).

Table 2 defines the cumulative S1–S4 input sets used for matched comparisons in Family 4. These labels refer to feature matrices, not model families. Raw composition is represented by eight atomic fractions when appropriate. Identifiable linear models use seven non-Ni fractions. The six Wen alloy descriptors are valence electron concentration (VEC), mixing enthalpy  $\Delta H_{\text{mix}}$ , mixing entropy  $\Delta S_{\text{mix}}$ ,  $\Omega$ , electronegativity mismatch  $\Delta\chi$ , and atomic-size mismatch  $\delta$  [12, 29, 30].




Figure 2: Representative candidate alloy verification: measured and target composition, backscattered-electron imaging, an EBSD inverse-pole-figure map, and the corresponding equivalent-circle grain-size distribution. The right-skewed distribution motivates retaining both mean grain size and its within-map standard deviation.

Table 2: Cumulative input sets used in Family 4.

| Set | $p$ | Content                                        |
|-----|-----|------------------------------------------------|
| S1  | 2   | $d^{-1/2}$ and $SD_{\text{grain}}$             |
| S2  | 8   | S1 plus six Wen alloy descriptors              |
| S3  | 11  | S2 plus cold work, $T_{rx}$ , and hold time    |
| S4  | 24  | S3 plus composition and SSS-related quantities |

### 3. Methods

#### 3.1. Family 1: Classical Hall-Petch

Equation (1) provides the common grain-size baseline. We also compare nine pre-specified laws:  $d^{-1/2}$ ,  $d^{-1}$ ,  $d^{-1/3}$ ,  $d^{-2/3}$ ,  $\ln(d)/d$ ,  $\ln d$ , two composite power-law forms, and a fitted free exponent. AIC, AICc, BIC, and Bayesian PSIS-LOO are used to compare these closely related parameterizations [31, 32, 33, 34]. We apply the same grain-size analysis separately to YS and HV.

#### 3.2. Family 2: Physics-derived descriptors

Family 2 tests whether compact physics-based representations explain YS and HV beyond the grain-size baseline. The Varvenne–Leyson–Curtin (VLC) model uses solute misfit volumes to estimate the zero temperature critical resolved shear

stress,


$$\tau_{y0} = 0.051\alpha^{-1/3}\bar{\mu}\left(\frac{1+\bar{\nu}}{1-\bar{\nu}}\right)^{4/3}f_1\left(\frac{\sum_i c_i\Delta V_i^2}{b^6}\right)^{2/3}, \quad (2)$$

followed by a thermal activation correction and a Taylor factor. Here,  $c_i$  and  $\Delta V_i$  are the elemental fractions and misfit volumes, and  $\bar{\mu}$ ,  $\bar{\nu}$ , and  $b$  are composition averaged elastic and lattice quantities. Weighted Labusch and Toda-Caraballo estimates provide two related SSS representations. The Labusch and Toda-Caraballo constants are calibrated once to the reported CoCrFeMnNi friction stress; VLC is left uncalibrated. The full expressions and parameter values are given in the Supplementary Information.

These implementations use Vegard law volumes and rule-of-mixtures elastic properties [8, 9, 10, 11]. We therefore keep calibrated and uncalibrated results separate. We also distinguish standalone prediction, improvement beyond Hall-Petch, and partial correlation after removing raw composition trends. These tests ask different questions, and partial correlation does not establish any form of independence.

The descriptor branch uses six Wen alloy descriptors together with  $d^{-1/2}$ ,  $SD_{\text{grain}}$ , and three processing variables. Standardization and six-component PCA are fit again within each training fold before OLS. Full-data loadings are descriptive only. This

**(a) Composition coverage by element and batch**



**(b) Processing coverage**




Figure 3: Composition and processing coverage by batch. The campaigns overlap but do not cover identical regions, and processing is partly structured by campaign. This data provenance motivates reporting LOBO along with record level validation instead of relying only on random splits.

fold-contained implementation corrects the preprocessing leakage in the earlier analysis.

### 3.3. Family 3: Composition/processing

Family 3 adds measured composition and processing to the Hall-Petch baseline through a pre-specified OLS hierarchy. Its principal model, M3, is

$$\sigma_y = \beta_0 + \sum_{e \neq \text{Ni}} \beta_e c_e + k_{\text{HP}} d^{-1/2}, \quad (3)$$

where Ni is the composition reference. Other models allow only V in the intercept, composition-dependent slope terms, variation in both intercept and slope, descriptor substitutions, or recorded processing variables. The family name includes processing because the broader hierarchy tests those covariates, although M3 itself uses composition and  $d$  without an explicit processing coefficient. Composition, processing, and grain size are interdependent in these campaigns. The fitted coefficients can therefore absorb covariance among them and are not interpreted as elemental constants.

The hierarchy is compared using information criteria and pooled out-of-fold predictions. We test  $\text{SD}_{\text{grain}}$  first as an additive microstructure term and then through M15,

$$\sigma_y = \beta_0 + \sum_{e \neq \text{Ni}} \beta_e c_e + k_{\text{HP}} d^{-1/2} + \gamma \text{SD}_{\text{grain}} + \eta \text{SD}_{\text{grain}} d^{-1/2}. \quad (4)$$

The interaction allows the grain-distribution width to modify the fitted grain-size contribution. It is pre-specified OLS and can be reproduced from the prepared dataset, but its coefficients remain associational because  $d$ ,  $\text{SD}_{\text{grain}}$ , composition, and processing are correlated.

### 3.4. Family 4: Non-linear estimators

Family 4 tests whether non-linear and ensemble estimators improve on linear controls when every model sees the identical S1–S4 matrices and validation splits, tuned and evaluated with ARMOTE-CV, which extends an automated regression workflow developed for other engineering domains [35, 36]. 20 estimators are compared: 15 (Bayesian ridge, ridge, Lasso, ElasticNet, kernel ridge, support vector regression (SVR), decision tree, random forest, extra trees, gradient boosting, XGBoost, LightGBM, CatBoost, Gaussian process regression (GPR), and a neural network) are tuned with nested multi-objective Optuna [37] search, and 5 (ordinary least squares (OLS), a multilayer perceptron (MLP), PCA-OLS (PCA = 6), a mean-predicting Dummy baseline, and a stacking ensemble) are run at fixed settings within this panel.

For every tuned model and every held-out split, hyperparameter optimization is a 5-fold CV carried out entirely on that split's training partition: candidate hyperparameters are scored across the 5 inner folds, the highest- $R^2$  configuration on the resulting MSE/ $R^2$  Pareto front is refit on the complete training partition,

and the refit model is then evaluated on the held-out validation partition, which contributes to neither the search nor the refit. Feature and target scalers are likewise fit on the training partition alone. The BIC is reported for the five linear-parametric estimators (ordinary least squares, ridge, Lasso, ElasticNet, and PCA-OLS) from the pooled LOO out-of-fold error, using the number of retained features plus one as the effective parameter count (six plus one for PCA-OLS). The full model roster, hyperparameter search spaces, and trial budget are given in the Supplementary Information.

### 3.5. Family 5: Symbolic regression

PySR searches three feature pools and three operator sets, while SISSO provides a sparse operator comparison [38, 39]. PySR pool P1 contains grain statistics, P2 adds raw composition and processing, and P3 adds Wen descriptors and processing. The operator sets progress from addition and subtraction to multiplication, division, and protected non-linear operators. SISSO starts from 27 composition, thermodynamic, SSS, and grain-size features and builds two algebraic expansion levels before sparse selection.

In the stored workflow, equations were selected after the complete data search fronts had been viewed. Their constants were then refit within the CV folds. We therefore call these results “post-selection fixed-form CV,” rather than unbiased outer-loop estimates of symbolic discovery. Protected square root and logarithm are evaluated as  $\sqrt{|x|}$  and  $\log(|x| + \epsilon)$ . We check each reported closed form for small or transformed denominators, sign changes, non-finite values, and extreme predictions over the observed internal and literature ranges. This audit is equation-specific and does not establish safety outside the stated descriptor envelope. BIC values from symbolic search are not compared directly with the pre-specified linear models because the stored symbolic parameter count omits the cost of searching variables, operators, and structures.

For HV, the selected structure is

$$H_V = c_0 + c_1 \frac{d_0 - d}{SD_{\text{grain}}} + c_2 \frac{\Delta H_{\text{mix}}}{t_{\text{hold}}^2}. \quad (5)$$

The constants are refit in every fold. In the reported equation,  $d$ ,  $d_0$ , and  $SD_{\text{grain}}$  are in  $\mu\text{m}$ ,  $t_{\text{hold}}$  is in h, and  $\Delta H_{\text{mix}}$  is in  $\text{kJ mol}^{-1}$ . Thus  $c_0$  and  $c_1$  have units of HV, while  $c_2$  has units of  $\text{HV} \cdot \text{h}^2 \cdot \text{mol kJ}^{-1}$ . This is an empirical expression, not a dimensionally universal relation. Its domain excludes  $SD_{\text{grain}} \rightarrow 0$  and  $t_{\text{hold}} \rightarrow 0$ .

### 3.6. Validation hierarchy and score uncertainty

We use three internal protocols that ask progressively different generalization questions: (i) shuffled 5-fold CV with seed 42; (ii) LOO; and (iii) LOBO across the six BO batches. LOO tests record-level interpolation, but another record with the same composition can remain in the training set. We therefore do not interpret LOO as strict new composition validation. LOBO tests transfer to a new iteration while retaining data from the broader experimental program and, in one case, the same composition.

For every protocol, predictions from all held-out records are pooled before computing

$$Q^2 = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2}. \quad (6)$$

We report this pooled quantity as cross-validated  $R^2$ . To calculate batch-clustered percentile intervals, we resample the six batches with replacement 1000 times while keeping every batch intact. These intervals show sensitivity to the sampled batches. With only six clusters, they should not be read as precise inferential confidence intervals. They quantify score uncertainty, not per-alloy predictive uncertainty; predictive intervals are outside the scope of this study.

The stored literature compilation has three evidence classes. E1 contains 54 direct strength records: 48 from the Borg/Citrine compilation [28] and six compression records for CrFeNi from Schneider and Laplanche [40]. These records represent only 18 exact compositions and mix sources and test modes. The legacy extraction limited grain size to 3–500  $\mu\text{m}$ , removed YS values above 800 MPa, and omitted one large-grain CrFeNi record because it departed from the expected Hall-Petch trend. Since two of these rules use the outcome or its expected relationship, we report E1 only as a heterogeneous literature stress test, not as an unbiased external validation set. E2 contains three values reconstructed from published Hall-Petch fits [4]. E3 contains 25 YS proxies converted from reported hardness [41]. We keep E2 and E3 separate from direct measurements.

For the hardness-strength comparison, hardness is converted to stress units and summarized as  $C_{\text{eff}} = \text{HV}(\text{MPa})/\sigma_y$ . Under the conditional Tabor-Hollomon approximation,

$$n_{\text{app}} = \frac{\ln(C_{\text{eff}}/3)}{\ln(0.08/0.002)} = \frac{\ln(C_{\text{eff}}/3)}{\ln 40}. \quad (7)$$

This is an apparent index, not a measured tensile hardening exponent, because offset YS is not the flow stress at exactly 0.002 true plastic strain and Hollomon behavior is assumed between 0.2% and 8% strain.

## 4. Results

Table 3 gives the common comparison across model families. Five-fold CV and LOO mainly test interpolation within the pooled dataset, whereas LOBO asks whether a model transfers to an unseen experimental batch. We therefore base the main conclusions on LOBO together with the change from record-level to batch-level validation, rather than on the highest score in any one column. A negative  $R^2$  means that the held-out predictions are worse than using the held-out mean.

The grouped scores should also be read in light of composition coverage. Six principal components represent 96.9% of the composition variance, and the mean pairwise batch-hull overlap in the PC1–PC2 projection is only about 0.21. CBC is relatively isolated in this projection, whereas CBA is well covered by the other batches. LOBO therefore combines easier interpolation-like folds with harder extrapolation folds; the pooled score is not a uniform measure of difficulty.

Table 3: Pooled out-of-fold  $R^2$  across the three internal validation protocols. LOBO provides the primary grouped transfer test. Family 4 includes linear controls to isolate the value of non-linearity. The Family 5 HV result is conditional on a structure selected using the full dataset.

| Target | Family/model          | Evidence status      | 5-fold | LOO   | LOBO   |
|--------|-----------------------|----------------------|--------|-------|--------|
| YS     | F1 Classical HP       | baseline             | 0.405  | 0.406 | 0.373  |
| YS     | F2 PCA-OLS            | fold-contained       | 0.462  | 0.480 | 0.362  |
| YS     | F3 M3                 | pre-specified linear | 0.666  | 0.652 | 0.625  |
| YS     | F3 M15 SD interaction | pre-specified linear | 0.731  | 0.694 | 0.694  |
| YS     | F4 Bayesian ridge S2  | nested               | 0.685  | 0.666 | 0.613  |
| YS     | F4 Kernel ridge S2    | nested               | 0.695  | 0.660 | 0.594  |
| HV     | F1 Classical HP       | baseline             | 0.086  | 0.136 | -0.077 |
| HV     | F4 Kernel ridge S1    | nested               | 0.228  | 0.466 | 0.190  |
| HV     | F5 fixed form         | post-selected        | 0.725  | 0.727 | 0.636  |

We do not report a single predictive ceiling. An in-sample composition-mean calculation gives  $R^2 \approx 0.86$  if all 72 singleton compositions in the 93-record YS subset are assigned zero residual. In contrast, a calculation based only on repeated compositions gives a formal value near  $-0.05$  because those groups also vary in processing and grain size. A separate estimate near 0.97 treats reported within-alloy standard deviations as irreducible target error, but replicate counts are unavailable. None of these quantities is a defensible out-of-sample ceiling.

#### 4.1. Family 1: Classical Hall-Petch

The classical Hall-Petch law provides the grain-size-only baseline. For YS, it gives training  $R^2 = 0.429$ , LOO  $R^2 = 0.406$  with RMSE = 62.5 MPa, and LOBO  $R^2 = 0.373$  (Fig. 4). Its similar LOO and LOBO scores show that mean grain size contains a modest, reproducible YS signal.

The exact grain-size law is less certain. Nine pre-specified forms were compared. Five two-parameter monotonic laws lie within  $\Delta\text{BIC} < 2$  and have LOO  $R^2$  from 0.395 to 0.406 (Fig. 5). The free-exponent fit gives  $n = 0.548$ , but its additional parameter raises  $\Delta\text{BIC}$  to 4.5. The Bayesian estimate is  $n = 0.53 \pm 0.16$ , with a 94% highest-density interval of about  $[0.22, 0.82]$ . Bayesian stacking places all weight on the classical form, but the candidate laws have nearly equal expected log predictive density, so this boundary weight is not decisive. The data support  $d^{-1/2}$  as a simple baseline, not as a uniquely identified exponent or mechanism.

The same baseline does not describe HV well. The classical form has training/LOO  $R^2 = 0.169/0.136$  and becomes negative under LOBO ( $-0.077$ ). A full-data optimized exponent of 1.728 raises LOO  $R^2$  only to 0.166 and is descriptive because the exponent was selected before CV. Mean grain size alone therefore does not transfer as a hardness model.

#### 4.2. Family 2: Physics-derived descriptors

Family 2 tests whether physics-derived descriptors add information beyond grain size and composition. For the SSS estimates, Cantor-anchored Labusch calibration gives a mean prediction-to-observation ratio of 1.07. Calibrated Toda-Caraballo and VLC overpredict by 42% and 78%, and uncalibrated variants overpredict by factors of 2–5. Used alone, VLC, Labusch, and Toda-Caraballo give LOO  $R^2 = 0.030, 0.274$ , and 0.091. Adding  $d^{-1/2}$  changes these values to 0.396, 0.469,




Figure 4: Family 1 Classical Hall-Petch fits for YS and HV. Grain size provides a useful baseline for YS, but mean grain size alone does not transfer for HV.




Figure 5: The three best-supported Family 1 grain-size laws fitted to YS over the observed range, with points coloured by batch. The curves overlap closely across the sampled grain sizes, which is why the data do not identify a unique exponent. The corresponding hardness comparison is given in the Supplementary Information.

and 0.399, compared with 0.406 for Hall-Petch alone (Table 4). Only calibrated Labusch gives a small improvement over that baseline (Fig. 6).

The descriptors do not add information once composition




Figure 6: Family 2 SSS predictions compared with Hall-Petch-corrected YS. Calibration can reduce average bias without showing that the descriptor transfers or adds information beyond composition.

is included. Adding an SSS estimate to a composition model changes LOO  $R^2$  by at most 0.002, and all linear partial correlations after removing composition are below 0.10 in magnitude. This redundancy applies to the present descriptor inputs; it does not distinguish limitations of SSS theory from uncertainty in atomic volumes, elastic properties, calibration, or unmeasured microstructure.

The broader Wen/PCA representation gives 5-fold/LOO/LOBO  $R^2 = 0.462/0.480/0.362$  when PCA is fitted inside each fold. Global preprocessing had produced 0.525/0.516/0.441, showing that even unsupervised scaling and PCA can inflate held-out scores when fitted before the split. The corrected LOBO result is below the classical Hall-Petch baseline, so the PCA model does not provide a transferable improvement at the batch level. Full-data loadings are therefore used only to describe the descriptor space, not as causal feature importance.

#### 4.3. Family 3: Composition/processing

Family 3 adds composition, processing, and measured microstructure terms to the Hall-Petch baseline. M3 gives LOO/LOBO  $R^2 = 0.652/0.625$  and LOBO RMSE = 49.7 MPa. Its batch-bootstrap interval for LOBO  $R^2$  is [0.501, 0.670], while the six individual batch-held-out scores range from 0.306 to 0.697. This spread shows that transfer difficulty differs across batches. M3 therefore retains useful YS information beyond mean grain size across the sampled experimental iterations.

The grain-size distribution adds information only in a specific form (Fig. 7). Adding  $SD_{\text{grain}}$  to M3 as an independent linear term raises LOO  $R^2$  from 0.652 to 0.668 but lowers LOBO from 0.625 to 0.595. M15, which includes both the additive term and  $SD_{\text{grain}}d^{-1/2}$ , gives 5-fold/LOO/LOBO  $R^2 = 0.731/0.694/0.694$  and LOBO RMSE = 44.9 MPa. Its LOO and LOBO batch-bootstrap intervals are [0.578, 0.766] and [0.599, 0.733]. Relative to M3, M15 lowers BIC by 20.1; the interaction coefficient has  $t = 4.71$  ( $p = 1.0 \times 10^{-5}$ ) in the complete-data OLS fit. M15 is therefore the strongest verified low-complexity YS model when  $SD_{\text{grain}}$  is measured. The additive control shows that the gain cannot be credited to grain-width measurement alone. Because the term is correlated with mean grain size and processing, it is interpreted as a predictive interaction, not as proof that distribution width changes the Hall-Petch slope.

The individual M3 LOBO scores are 0.697 for BBA, 0.306 for BBB, 0.656 for BBC, 0.423 for CBA, 0.584 for CBB, and

0.409 for CBC; the corresponding RMSE values range from 32.9 to 60.4 MPa. CBC has the lowest projected hull containment (0.476), while CBA has complete containment in the PC1-PC2 projection. Hull containment is one of several factors behind the batch-level variation. These folds show why the pooled LOBO score should be accompanied by batch-level results.

The complete-data M3 fit gives  $k_{\text{HP}} = 766 \text{ MPa} \cdot \mu\text{m}^{1/2}$ . Adding the composition-dependent intercept lowers the slope from about  $1182 \text{ MPa} \cdot \mu\text{m}^{1/2}$  in the grain-only fit, consistent with composition-grain-size confounding in Family 1. Reported room-temperature values for concentrated FCC alloys span 265 to  $1244 \text{ MPa} \cdot \mu\text{m}^{1/2}$  with a median of 770 [5, 4, 42, 40, 6, 7], so the fitted slope is close to the middle of the published set. The intercepts fall in the published range as well. Across the 94 conditions M3 gives  $\sigma_0$  between 84 and 269 MPa with a mean of 151, against 80 to 243 MPa and a median of 149 in the same studies. The range of reported slopes is wide partly because studies of the same alloy disagree: CoCrFeNi has been reported at 276, 506, 684, and  $855 \text{ MPa} \cdot \mu\text{m}^{1/2}$  [42, 6, 7, 43]. Much of that spread tracks the measurement rather than the material. The coefficient depends on test temperature, falling from 538 to  $127 \text{ MPa} \cdot \mu\text{m}^{1/2}$  between 77 and 1073 K for a single alloy [4]. It depends on whether annealing twins are counted as boundaries, with the same CrFeNi specimens giving  $966 \text{ MPa} \cdot \mu\text{m}^{1/2}$  against grain size and 897 against crystallite size [40]. It depends on the window of grain sizes fitted, and the two standard pure-metal surveys report 110 and  $241 \text{ MPa} \cdot \mu\text{m}^{1/2}$  for copper from data sets reaching different lower limits [3, 44]. A coefficient fitted to converted hardness is a different quantity again: one review reports  $226 \text{ MPa} \cdot \mu\text{m}^{1/2}$  for CoCrFeMnNi after dividing published microhardness by a Tabor factor of three, against 494 measured in tension for the same alloy [45, 4]. Converted coefficients are therefore excluded from Table ???. A fitted slope is comparable to a literature value only when temperature, test mode, boundary-counting convention, and grain-size window all match, and the M3 value should not be treated as a universal FCC value. Among the composition coefficients (Table 5), V is the largest positive term relative to Ni (+291 MPa), but all coefficients are conditional on the sampled compositions and processing routes.

Two sensitivity checks address numerical stability of this fit. Across 1000 grain-size perturbations, the V, Co, Cr, and Fe coefficients change by less than about 15%, whereas  $k_{\text{HP}}$  changes by about 39%. Al is not included in the 15% statement because it spans only 0–4 at.% and has a much wider interval. A 10,000-replicate row bootstrap gives standard-error ratios of 0.99–1.54 relative to OLS. These checks support stability within the sampled records, but a row bootstrap does not reproduce campaign-level sampling and cannot establish transfer.

M3 assumes a shared Hall-Petch slope, but the dataset cannot test that assumption decisively. Only nine compositions are repeated, most have two records, and processing changes within some groups. Their descriptive local slopes range from about  $-8539$  to  $9282 \text{ MPa} \cdot \mu\text{m}^{1/2}$ . A fixed-effect summary gives a slope near 2144, Cochran  $Q \approx 251$ , and  $I^2 = 96.8\%$ , but this apparent heterogeneity cannot be assigned to composition under the present design. The data therefore do not show that  $k_{\text{HP}}$  is composition-independent. The supported result is narrower: a

Table 4: Family 2 YS descriptor results. Ratio is mean SSS prediction divided by mean Hall–Petch-corrected YS. Raw  $R^2$  and Spearman  $\rho$  compare the complete-data SSS estimates with the corrected YS. Standalone and +HP are LOO results; +HP is Ridge using the SSS estimate and  $d^{-1/2}$ .

| Model                      | Mean (MPa) | Ratio | Raw $R^2$ | Spearman $\rho$ | Standalone $R^2$ | +HP $R^2$ | +HP RMSE (MPa) |
|----------------------------|------------|-------|-----------|-----------------|------------------|-----------|----------------|
| VLC                        | 342        | 1.78  | −6.92     | 0.53            | 0.030            | 0.396     | 63.1           |
| Labusch, calibrated        | 206        | 1.07  | 0.14      | 0.55            | 0.274            | 0.469     | 59.1           |
| Toda–Caraballo, calibrated | 274        | 1.42  | −3.18     | 0.49            | 0.091            | 0.399     | 62.9           |
| HP only                    | —          | —     | —         | —               | —                | 0.406     | 62.5           |

Table 5: Family 3 M3 coefficients relative to Ni. Intervals are conditional on the OLS specification and sampled campaigns.

| Parameter | Estimate (MPa) | 95% low | 95% high |
|-----------|----------------|---------|----------|
| Intercept | 229.6          | 152.4   | 306.7    |
| Al        | −201.3         | −955.4  | 552.8    |
| Co        | −187.2         | −268.1  | −106.3   |
| Cr        | −307.7         | −588.5  | −26.9    |
| Cu        | −334.0         | −651.9  | −16.2    |
| Fe        | −360.4         | −520.1  | −200.6   |
| Mn        | 82.3           | −117.6  | 282.2    |
| V         | 291.3          | 41.0    | 541.5    |
| $k_{HP}$  | 765.8          | 501.5   | 1030.0   |

composition-dependent intercept with a shared grain-size term predicts YS well across the sampled batches. The physical division among friction stress, processing, and grain-boundary strengthening remains unresolved.

#### 4.4. Family 4: Non-linear estimators

Family 4 tests whether non-linear and ensemble estimators improve on linear controls when every model sees the identical, matched S1–S4 inputs under nested tuning. 5-fold CV and LOO mainly probe interpolation within the pooled 93/94-record dataset; LOBO asks whether a model transfers to an unseen experimental batch. Reading all three together, rather than any one score in isolation, is what exposes whether a model’s apparent skill survives batch-level transfer. Full 5-fold and LOO scores for all fifteen reported estimators across S1–S4 are given in the Supplementary Information (SI).

For YS, regularized linear and kernel methods lead under every validation protocol, not just LOBO. At S2, ridge gives 5-fold/LOO/LOBO  $R^2 = 0.695/0.662/0.614$  and Bayesian ridge gives 0.685/0.666/0.613, essentially indistinguishable from one another across all three metrics. This score is directly verifiable: reloading the stored per-fold models and scalars and re-predicting each held-out batch reproduces it exactly, with pooled RMSE = 50.5 MPa. Kernel ridge is the closest non-linear competitor (0.695/0.660/0.594); tree-based and boosted estimators trail further behind under LOBO specifically (LightGBM 0.514/0.565/0.555, XGBoost 0.564/0.561/0.527, random forest 0.588/0.563/0.498), even where their 5-fold or LOO scores are comparable to the linear leaders. The gap between sample-level and batch-level performance is accordingly smaller for the linear and kernel models (5-fold-to-LOBO drop of 0.07–0.10) than for several tree ensembles, whose apparently competitive sample-level scores decline more under LOBO. The same ordering, linear and kernel methods ahead of tree ensembles under LOBO, holds at S1, S3, and S4 (Fig. 8).

5 estimators are omitted from Fig. 8 and this comparison for two distinct reasons. Lasso is excluded because its nested search collapses it to an intercept-only fit at S2 and S3: its LOBO  $R^2$  (−0.122) is identical to the Dummy mean-predicting baseline at both tiers, and its LOO score is likewise close to Dummy’s. Lasso’s tuned  $\alpha$  range, applied to a standardized target, drives all coefficients to exactly zero at S2 and S3, collapsing it to an intercept-only fit. Ridge, Bayesian ridge, and ElasticNet do not show this behavior since their L2 or blended L1/L2 penalties shrink coefficients rather than zeroing them outright. With random forest, extra trees, gradient boosting, XGBoost, and LightGBM already spanning the tree-ensemble and boosting space, CatBoost and decision tree add no distinct performance improvement or additional information. The MLP is excluded for instability rather than redundancy as its scores swing between competitive and severely negative across tiers and protocols (YS S1 LOBO = −1.03; HV S3 LOBO = −2.28). NNR, the TensorFlow implementation with tunable neural architecture, shows similar volatility but is retained as MLP’s fixed configuration reflects the instability in its architecture rather than a fold-specific search failure.

Set against Family 1–3, Family 4’s most flexible search does not exceed what a pre-specified, low-dimensional linear model already achieves. The classical Hall–Petch baseline (Family 1) gives LOBO 0.373; the PCA-OLS descriptor model (Family 2) gives 0.362; the composition Hall–Petch model M3 (Family 3) gives 0.625; and the  $SD_{\text{grain}}$ -interaction model M15 (Family 3) gives 0.694. Family 4’s best result, ridge or Bayesian ridge at S2, reaches only 0.613–0.614, below both Family 3 models despite drawing on twenty estimators and a fold-nested hyperparameter search. M15 is also more stable across protocols than anything Family 4 produces. Its 5-fold-to-LOBO gap is 0.037, smaller than ridge’s 0.081 or kernel ridge’s 0.101. A pre-specified, physically motivated interaction term generalizes further, and more consistently across validation regimes, than an automated search over a much larger model space.

HV shows the same three-protocol logic delivering a starker result. At S1, kernel ridge is the only estimator among all twenty with a positive LOBO score (0.228/0.466/0.190); every other linear, kernel, tree, and ensemble method is negative under LOBO at this tier. Once composition or processing descriptors are added (S2–S4), the sample-level and batch-level pictures diverge sharply. Several estimators still post substantial LOO scores at S2, for example LightGBM (LOO = 0.520) and kernel ridge (LOO = 0.449), while their LOBO scores collapse to −0.911 and −0.559, respectively. Every estimator we tested at S2, S3, and S4 scores at or below the Dummy regressor base-

Family 3 Hall-Petch hierarchy for yield strength. Negative  $\Delta\text{BIC}$  favours the model over M3.




Figure 7: Family 3 Hall-Petch hierarchy for YS. M15 adds  $\text{SD}_{\text{grain}}$  to M3 both directly and through the interaction  $\text{SD}_{\text{grain}}d^{-1/2}$ . It gives the strongest internal result, but its coefficients remain conditional on the sampled campaigns and correlated microstructure.




Figure 8: Family 4 matched-input LOBO comparison across the S1-S4 ladder for YS and HV. 15 of 20 estimators are shown; Lasso, CatBoost, decision tree, the multilayer perceptron, and the Dummy baseline are omitted. Regularized linear and kernel-ridge estimators lead for YS at every tier. For HV, only kernel ridge is positive at S1; no estimator beats the held-out mean once composition or processing descriptors are added (S2-S4).

line under LOBO; the least-negative score at each of those tiers belongs to Dummy itself. LOO alone would suggest several usable HV models past S1 but LOBO shows none of them transfer to an unseen batch. Compared with Family 1 (classical HP, LOBO = -0.077), Family 4 HV illustrates a wide sample-level-to-batch-level gap in LOBO. No estimator, of any flexibility, transfers past grain size alone.

#### 4.5. Family 5: Symbolic regression

Family 5 searches for compact analytical forms. For YS, the two accuracy-selected PySR equations using composition

and processing each reach fixed-form LOO  $R^2 = 0.808$  after their constants are refitted. Their archived mean per-batch LOBO scores are 0.696 and 0.689. A descriptor-based accuracy equation gives LOO/stored-LOBO  $R^2 = 0.765/0.668$ , while the descriptor-based elbow equation gives 0.674/0.555. These archived LOBO values are means of fold-wise  $R^2$  and are not directly comparable with the pooled  $Q^2$  values in Table 3. The descriptor-based accuracy equation contains  $\text{SD}_{\text{grain}} / \sqrt{|\Delta S_{\text{mix}} - 8.2739|}$ . The internal minimum  $\Delta S_{\text{mix}} = 8.3637$  is only 0.0898 above its pole. The protected square root keeps the expression real on either side, but does not remove the denominator singularity. This equation is therefore unsafe outside the observed descriptor envelope.

All of these PySR values are post-selection scores: the equation structures were chosen from search fronts built with the full dataset. They are not nested estimates of symbolic-regression performance and cannot be ranked directly against the pre-specified models. The complexity-loss front from which they were taken is shown in Fig. 9. PySR BIC also counts fitted constants but not the cost of searching variables, operators, and structures, so it is not compared with the BIC of the pre-specified linear hierarchy. The SISSO equations have the same selection limitation. Grain-width terms recur across searches and are useful hypotheses, but not independent evidence of a mechanism: all searches use the same 94 records, and  $d$  is correlated with  $\text{SD}_{\text{grain}}$ .

Refitting the selected HV structure gives

$$H_V = 222.3 + 84.5 \frac{6.93 \mu\text{m} - d}{\text{SD}_{\text{grain}}} - 1.005 \frac{\Delta H_{\text{mix}}}{t_{\text{hold}}^2}. \quad (8)$$

The three coefficient bootstrap intervals are [202.9, 240.1], [69.4, 98.3], and [-1.218, -0.803]. The corresponding bootstrap




Figure 9: Family 5 PySR complexity-loss front. Choosing an elbow or accuracy equation after viewing this front places structure selection outside the later constant-refit CV, so the result must be labeled post-selection.

means and standard deviations are  $221.82 \pm 9.64$ ,  $84.12 \pm 7.34$ , and  $-1.006 \pm 0.106$ ; each interval excludes zero, conditional on the selected structure. The full-data fit gives  $R^2 = 0.744$  and  $\text{RMSE} = 16.90$  HV. With the structure fixed and only its constants refitted, 5-fold/LOO/LOBO  $R^2 = 0.725/0.727/0.636$ , with  $\text{RMSE} = 17.50/17.46/20.16$  HV. The batch-bootstrap interval for LOBO  $R^2$  is  $[0.329, 0.765]$ .

The equation retains useful batch-held-out performance, but its structure was selected after viewing the full-data search fronts. The zero at  $d = 6.93 \mu\text{m}$  is not a singularity; the unsafe limits are  $\text{SD}_{\text{grain}} = 0$  and  $t_{\text{hold}} = 0$ . The observed data remain away from both limits. The two non-constant terms have different provenance risks. Hold time is fixed at 0.5 h for all 35 BB-campaign records and takes values of 1, 2, 4, or 8 h in Campaign CB, so hold time separates campaign with  $\text{AUC} = 1.000$  and cannot be identified within Campaign BB. The combined  $\Delta H_{\text{mix}}/t_{\text{hold}}^2$  term also separates campaign strongly ( $\text{AUC} = 0.970$ ). In contrast,  $(6.93 - d)/\text{SD}_{\text{grain}}$  is nearly campaign-neutral ( $\text{AUC} = 0.533$ ). The hold-time term must therefore be treated as a campaign proxy in these data; the grain-width term remains a post-selected, correlated microstructure hypothesis. Testing either term as a mechanism requires a prospectively frozen equation and direct distribution descriptors such as  $\text{SD}/d$ , lognormal width, or  $E[d^{-1/2}]$ .

The SISSO comparison also separates numerical safety from predictive accuracy. SISSO Full places  $\delta_\mu$  in a denominator and becomes unstable when constituent shear moduli are similar. Removing that term lowers the aggregate literature RMSE from 420.6 to 162.9 MPa. The bounded form is numerically safer, but its aggregate  $R^2 = -1.37$  is still unsuitable for prediction.

Additional symbolic controls give lower or unstable results. An expanded SISSO search reports LOO/LOBO  $R^2 =$




Figure 10: HV rank versus YS rank by batch. Strong within-batch ordering becomes weaker after batches are pooled because their composition and processing coverage differs.

0.604/0.350 when the selected transforms change across folds. A universal-operator EML control gives LOO  $R^2 = 0.321$  at depth one and becomes unstable at greater depth. A recalibrated pure-metal descriptor relation gives LOO  $R^2 = 0.346$ . These controls suggest that restricted operators can provide useful inductive bias, but the archived files do not establish that every structure-selection step was nested.

#### 4.6. Hardness-yield correspondence

The different YS and HV results motivate a direct comparison of the two properties. Across the 93 paired records,  $C_{\text{eff}} = 5.13 \pm 1.36$ , corresponding through Eq. (7) to  $n_{\text{app}} \approx 0.15 \pm 0.07$  after propagating the reported  $C_{\text{eff}}$  summary. The global Spearman correlation is 0.46, although the within-batch correlations range from 0.70 to 0.95. Pooling weakens, but does not reverse, the association; this is aggregation-induced attenuation rather than a strict Simpson paradox. The partial rank correlation after controlling for  $d$  is only 0.24.

The archived descriptive relation is  $\log C_{\text{eff}} = 1.36 + 0.10 \log d - 2.06 x_V$ , with  $d$  entered in the fitted numerical units. It is not dimensionally invariant and has not been prospectively validated. Ranking agreement also depends on the cutoff (Fig. 10): 2 of the 5 hardest alloys are among the 5 strongest, increasing to 8/10 and 11/20 for broader sets. HV may help rank alloys within a narrow composition and processing window, but the results do not support a single transferable HV-to-YS conversion or identify V as a causal source of the pooled attenuation.

#### 4.7. Literature stress test

The E1 literature set provides a final stress test outside the experimental program. On its 54 direct-strength records, M3 has the lowest RMSE, 149.5 MPa, but  $R^2 = -0.508$ . SISSO Robust gives  $\text{RMSE} = 160.0$  MPa and  $R^2 = -0.726$ , while

Table 6: Legacy E1 direct-strength literature stress test. Negative  $R^2$  indicates performance worse than predicting the E1 mean.

| Model           | $R^2$  | RMSE (MPa) | Bias (MPa) |
|-----------------|--------|------------|------------|
| F3 M3           | -0.508 | 149.5      | -76.6      |
| F5 SISSO Robust | -0.726 | 160.0      | -61.3      |
| F5 SISSO Full   | -5.084 | 300.3      | +50.0      |

the small-denominator term in SISSO Full contributes to  $R^2 = -5.084$ . All three models are worse than predicting the E1 mean (Table 6). The test is useful for exposing domain sensitivity and unsafe equations, but not for estimating unbiased external performance: E1 mixes sources and test modes and was filtered using outcomes.

The other literature classes give a different appearance because they are not independent direct measurements. On the three E2 values reconstructed from published Hall–Petch fits, M3 and SISSO Robust give descriptive  $R^2 = 0.677$  and  $0.849$ , with RMSE = 49.0 and 33.5 MPa. With  $n = 3$ , these  $R^2$  values are not interpretable as generalization estimates. On the 25 E3 hardness-converted proxies, the same models give  $R^2 = -2.157$  and  $-9.642$ . Pooling E1–E3 gives M3, SISSO Robust, and SISSO Full  $R^2 = -0.580$ ,  $-1.373$ , and  $-14.813$ , respectively. Across all 18 stored PySR YS variants, the pooled E1–E3  $R^2$  is also negative; the best is  $-0.031$  (RMSE = 107.4 MPa), and the highlighted descriptor accuracy equation gives  $-0.295$  (RMSE = 120.4 MPa). These pooled values are reported only as numerical stress tests; they do not increase the amount of direct external evidence.

## 5. Discussion

### 5.1. Yield strength across the five families

The YS results form a clear progression. Family 1 establishes a useful grain-size baseline, but several nearby exponents fit almost equally well. Family 2 shows that the tested SSS and Wen/PCA representations do not provide a transferable improvement with the available approximations. Family 3 supplies the largest verified gain: M3 combines a composition-dependent intercept with a shared Hall–Petch slope, while M15 adds a grain-width interaction and improves LOBO from 0.625 to 0.694. The additive  $SD_{\text{grain}}$  control does not transfer as well, so the result is specific to the interaction. Family 4 reaches similar batch-held-out accuracy with a linear control and a non-linear model, so estimator flexibility is not the source of the gain. Family 5 produces compact equations and further grain-width hypotheses, but its equation search was not nested and its literature performance is negative. M15 is the strongest verified low-complexity YS model when grain-distribution width is measured; M3 remains the simpler choice when only mean grain size is available. Neither is a universal Hall–Petch law.

### 5.2. Hardness requires a different model

The same sequence gives a different result for HV. Mean grain size, physics descriptors, composition, and the tested non-linear models all transfer poorly across batches. Adding  $SD_{\text{grain}}$

linearly improves LOO but worsens LOBO. The post-selected Family 5 fixed form performs well under LOBO after coefficient refitting, but its structure selection was not nested. Its grain-width term is nearly campaign-neutral, whereas its hold-time term closely tracks campaign identity. These terms are worth testing separately, but they do not yet define a hardness law. The variable  $C_{\text{eff}}$  and the rank analysis also show why HV should not be treated as a direct substitute for YS across different composition and processing domains.

### 5.3. The scientific conclusion depends on validation scale

The same fitted model can support different conclusions at different validation scales. The clearest examples are for HV: classical Hall–Petch changes from LOO  $R^2 = 0.136$  to LOBO  $R^2 = -0.077$ , LightGBM S1 falls from 0.323 to 0.112, and the legacy tuned XGBoost model falls from 0.729 to 0.574. In Family 2, global PCA gives LOO/LOBO  $R^2 = 0.516/0.441$ , whereas fold-contained PCA gives 0.480/0.362, showing how preprocessing and grouping compound. In contrast, M3 changes only modestly from 0.652 to 0.625, and M15 gives 0.694 under both protocols. The Family 5 HV fixed form retains  $R^2 = 0.636$  across batches, but this estimate remains conditional on a structure selected from the full dataset and includes a hold-time term confounded with campaign. LOBO is a useful internal transfer test, but it is not equivalent to deployment because the batches share program-level synthesis, characterization, and design practices.

M15 is therefore the best verified YS model for record-level interpolation and new-batch prediction within the combined experimental program when  $SD_{\text{grain}}$  is available. M3 is a simpler alternative that does not require grain-distribution width. Neither model is a universal Hall–Petch law or a strict test on unseen compositions. Lasso is also useful as a matched predictive control, although it has less direct physical structure. The Family 5 equations remain discovery hypotheses because their structure selection was not nested and their literature evidence is unfavorable. For HV, Eq. (8) is the strongest candidate within the sampled batches, but it should be frozen and tested prospectively before it is treated as a transferable hardness law.

### 5.4. Physical interpretation without overreach

The data support a grain-size baseline for YS together with additional composition- and grain-distribution-associated signals. They do not uniquely attribute those signals to friction stress, solute misfit, boundary chemistry, processing, or hetero-deformation-induced strengthening. The apparent preference for a constant Hall–Petch slope may result from limited identifiability rather than true invariance.  $SD_{\text{grain}}$  appears in M15, the Family 4 feature sets, and several symbolic expressions, but this recurrence is not independent confirmation because the models use the same records and correlated predictors. Raw standard deviation scales with the mean and can also depend on mapped area, grain count, weighting, and segmentation.

YS and HV are not interchangeable labels. Across the 93 paired records, the effective ratio  $HV(\text{MPa})/\sigma_y$  is  $5.13 \pm 1.36$ . Indentation imposes roughly 8% plastic strain, so in an alloy that

work hardens the indentation samples a flow stress well above the offset yield point and the ratio has to exceed three [13, 44]. The published records behave the same way. Among the compiled FCC alloys that report grain size, yield strength and hardness for the same material, the ratio runs from 4.2 to 6.3 [28]. Across the wider set that reports strength and hardness without grain size it runs from 3.2 to 17.6, with the lowest values in heavily deformed samples that have little work-hardening capacity left. Strong within-batch rank correlations become weaker when the campaigns are pooled. A single Tabor conversion cannot remove the effects of work hardening, constraint, processing, and campaign structure. Separate models and validation are needed for each property, even though both measurements arise from related plastic-deformation processes.

### 5.5. Best practices for physics-informed property ML

The five families are most useful as a sequence of scientific questions, not as a leaderboard. Family 1 defines a physically interpretable baseline. Family 2 tests whether physics-based descriptors explain the remaining error. Family 3 asks what measured composition and processing add. Family 4 tests estimator flexibility using matched inputs and budgets. Family 5 generates closed-form hypotheses and therefore requires the strictest audit of selection and domain.

For all five families, the validation split must match the proposed use. LOO measures record-level interpolation and must be interpreted alongside the repeated-composition audit. LOBO is needed for claims about a new experimental iteration. Preprocessing, feature selection, hyperparameter tuning, and equation discovery must all occur within that validation boundary. Uncertainty estimates should preserve experimental clusters. Direct measurements should remain separate from values reconstructed from equations or converted from other properties. External sets should be defined before outcomes are inspected and should document composition, processing, grain statistics, phase state, temperature, and test mode. Closed forms also require a singularity audit over both observed and intended deployment domains.

Table 7 maps each claim made in this work onto the evidence that supports it and the interpretation that evidence allows. Mechanistic claims require more than agreement among predictive models. A proposed term should survive alternative parameterizations, dimensional analysis, propagation of measurement uncertainty, and a prospective experiment that varies the proposed cause independently of confounders. For the present grain-width hypothesis, such an experiment should hold mean grain size constant while changing the distribution shape. It should also record complete EBSD distributions and sampling metadata, with the equation frozen before testing.

### 5.6. Limitations and reproducibility

The analysis includes six batches from two sequential campaigns within one research program. Bayesian-optimization sampling narrows the observed domain, and the composition and processing coverage differs among batches. Exact replicates are limited, processing and composition are not orthogonal, and complete grain-size distributions are reduced to a mean and standard deviation. The literature stress set is not suitable for model

selection or an external accuracy claim. A future benchmark should be registered before prediction and should use consistent tensile definitions and complete processing metadata.

The revised repository separates reproducible results from archival, post-selected, and proxy-based results. The scripts `grouped_validation.py` and `sdgrain_models.py` generate the primary grouped-validation table and the M3/M15 grain-width audit from the prepared descriptor dataset. The standard OLS, information-criterion, and grouped-CV results for M15 are independently reproduced here. Its earlier Bayesian stacking weight and PSIS-LOO comparison remain archival because that probabilistic analysis was not regenerated. The ARMOTE-CV panel stores a separate Optuna study, parameter set, and scaler for every fold, and its reported scores are reproducible from those objects, so its nested implementation is verified. PySR equations were instead selected on complete-data Pareto fronts; we refit their constants per fold but did not rediscover their structure inside a training fold. Both would require fold-level generators and stored out-of-fold predictions to verify, and we therefore treat both as archival.

## 6. Conclusions

Physics-informed ML improves materials-property modeling only when physics constrains both the representation and the claim. In the fixed hierarchy of Fig. 1, Family 1 provides the grain-size baseline. Family 2 shows that the tested SSS and Wen/PCA descriptors are redundant or conditional on the campaign. Family 3 gives the clearest verified improvement for YS through a composition-dependent Hall–Petch intercept and a grain-width interaction. Family 4 shows no demonstrated advantage for non-linearity when the inputs are matched. Family 5 produces compact hypotheses, but their selection and transfer remain unresolved.

The key result is the gap between record-level and batch-level validation. M15 gives the strongest verified low-complexity YS performance when  $SD_{\text{grain}}$  is measured, while M3 retains moderate performance without that feature. The tested non-linear estimator does not improve on its linear control at matched inputs. The symbolic HV form performs best across the sampled batches but remains conditional on post-selected structure and a hold-time term confounded with campaign. The practical implication is straightforward. Start with the known physics, test compact descriptors, then add measured composition, processing, and microstructure before increasing estimator complexity. Choose validation groups that match the intended decision. A model should not progress from interpolation to physical interpretation or alloy design without fold-contained selection, cluster-preserving score intervals, outcome-independent external data, and prospective mechanistic tests.

### CRedit authorship contribution statement

**Mrinalini Mulukutla:** Methodology, Software, Investigation, Formal analysis, Data curation, Visualization, Project administration, Writing – original draft. **Shakti Prasad Padhy:**

Table 7: Claim–evidence map after grouped-validation and provenance audits.

| Claim                                  | Evidence                                                                                   | Supported interpretation                                                            |
|----------------------------------------|--------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|
| Composition improves YS                | M3 LOO/LOBO $R^2 = 0.652/0.625$                                                            | Supported for record- and batch-level prediction within the present program         |
| Grain-width interaction improves YS    | Additive SD gives 0.668/0.595; M15 gives 0.694/0.694                                       | Supported as a predictive interaction when SD is measured; mechanism not identified |
| Non-linearity improves YS              | Lasso/LightGBM LOBO 0.621/0.615                                                            | No demonstrated advantage for the tested non-linear model at matched S2 inputs      |
| The HV closed form generalizes         | Fixed-form LOO/LOBO 0.727/0.636; post-selected structure                                   | Strongest internal HV candidate; prospective validation is still required           |
| Grain-width terms identify a mechanism | Shared data, correlated $d$ and $SD_{\text{grain}}$ , post-selected forms                  | Testable hypothesis requiring controlled microstructure and prospective validation  |
| Models deploy to literature alloys     | All evaluated E1 stress-test $R^2 < 0$ ; set has mixed modes and outcome-informed curation | No deployment claim and no unbiased external estimate                               |

Methodology, Software, Formal analysis, Project administration, Visualization, Writing – original draft. **Wenle Xu:** Data curation, Investigation, Methodology, Writing – original draft. **Bibhu Prasad Sahu:** Data curation, Investigation, Methodology, Writing – original draft. **Joydeep Kundu:** Data curation, Investigation, Methodology, Writing – original draft. **Vasanth C. Shunmugasamy:** Data curation, Investigation, Methodology, Writing – original draft. **Douglas Allaire:** Conceptualization, Formal analysis, Funding acquisition, Software, Supervision, Writing – review & editing. **Ibrahim Karaman:** Conceptualization, Funding acquisition, Supervision, Writing – review & editing. **Raymundo Arróyave:** Conceptualization, Formal analysis, Funding acquisition, Software, Methodology, Supervision, Writing – review & editing.

### Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

### Acknowledgments

This research was sponsored by the Army Research Laboratory under Cooperative Agreement Number W911NF-22-2-0106. The views and conclusions in this document are those of the authors and should not be interpreted as official policies of the Army Research Laboratory or the US Government. The US Government is authorized to reproduce and distribute reprints for Government Purposes, notwithstanding any copyright notation herein.

The authors would like to acknowledge Dr. Anup Bandyopadhyay at Texas A&M University for access to the XRD Bruker D8 Discover instrument, and Texas A&M’s Materials Characterization Facility for access to their EBSD-capable FIB-SEM. We also acknowledge Texas A&M’s High Performance Research Computing group for access to the computational infrastructure used in this work.

### Data and code availability

The data generated and analyzed during this study are included in the Supplementary Information and are also available as an Excel database in the GitHub repository below.

All code used in this project, including the ARMOTE-CV non-linear ML modeling, is available at <https://github.com/ArroyaveLab/2026-mulukutla-hall-petch>.

### Declaration of generative AI and AI-assisted technologies in the writing process

During the preparation of this work, the authors used Claude Opus 5 and ChatGPT 5.6 Sol for data analysis assistance and prose editing. After using these tools, the authors reviewed and edited the content as needed and take full responsibility for the content of the publication.

### References

- [1] E. O. Hall, The deformation and ageing of mild steel: III. Discussion of results, *Proceedings of the Physical Society. Section B* 64 (9) (1951) 747–753. doi:10.1088/0370-1301/64/9/303.
- [2] N. J. Petch, The cleavage strength of polycrystals, *Journal of the Iron and Steel Institute* 174 (1953) 25–28.
- [3] Z. C. Cordero, B. E. Knight, C. A. Schuh, Six decades of the Hall–Petch effect – a survey of grain-size strengthening studies on pure metals, *International Materials Reviews* 61 (8) (2016) 495–512. doi:10.1080/09506608.2016.1191808.
- [4] F. Otto, A. Dlouhý, C. Somsen, H. Bei, G. Eggeler, E. P. George, The influences of temperature and microstructure on the tensile properties of a CoCrFeMnNi high-entropy alloy, *Acta Materialia* 61 (15) (2013) 5743–5755. doi:10.1016/j.actamat.2013.06.018.

- [5] S. Yoshida, T. Bhattacharjee, Y. Bai, N. Tsuji, Friction stress and Hall–Petch relationship in CoCrNi equiatomic medium entropy alloy processed by severe plastic deformation and subsequent annealing, *Scripta Materialia* 134 (2017) 33–36. doi:10.1016/j.scriptamat.2017.02.042.
- [6] J. Li, K. Yamanaka, Y. Zhang, T. Furuhashi, G. Cao, J. Hu, A. Chiba, Reaching unconventionally large Hall–Petch coefficients in face-centered cubic high-entropy alloys, *Materials Research Letters* 12 (6) (2024) 399–407. doi:10.1080/21663831.2024.2337211.
- [7] A. Jagetia, M. S. K. K. Y. Nartu, S. Dasari, A. Sharma, B. Gwalani, R. Banerjee, Ordering-mediated local nano-clustering results in unusually large Hall–Petch strengthening coefficients in high entropy alloys, *Materials Research Letters* 9 (5) (2021) 213–222. doi:10.1080/21663831.2020.1871440.
- [8] C. Varvenne, A. Luque, W. A. Curtin, Theory of strengthening in fcc high entropy alloys, *Acta Materialia* 118 (2016) 164–176. doi:10.1016/j.actamat.2016.07.040.
- [9] C. Varvenne, W. A. Curtin, Predicting yield strengths of noble metal high entropy alloys, *Scripta Materialia* 142 (2017) 92–95. doi:10.1016/j.scriptamat.2017.08.030.
- [10] R. Labusch, A statistical theory of solid solution hardening, *Physica Status Solidi (b)* 41 (2) (1970) 659–669. doi:10.1002/pssb.19700410221.
- [11] I. Toda-Caraballo, P. E. J. Rivera-Díaz-del Castillo, Modelling solid solution hardening in high entropy alloys, *Acta Materialia* 85 (2015) 14–23. doi:10.1016/j.actamat.2014.11.014.
- [12] C. Wen, C. Wang, Y. Zhang, S. Antonov, D. Xue, T. Lookman, Y. Su, Modeling solid solution strengthening in high entropy alloys using machine learning, *Acta Materialia* 212 (2021) 116917. doi:10.1016/j.actamat.2021.116917.
- [13] D. Tabor, *The Hardness of Metals*, Clarendon Press, Oxford, 1951.
- [14] M. Dao, N. Chollacoop, K. J. Van Vliet, T. A. Venkatesh, S. Suresh, Computational modeling of the forward and reverse problems in instrumented sharp indentation, *Acta Materialia* 49 (19) (2001) 3899–3918. doi:10.1016/S1359-6454(01)00295-6.
- [15] B. Cantor, I. T. H. Chang, P. Knight, A. J. B. Vincent, Microstructural development in equiatomic multicomponent alloys, *Materials Science and Engineering: A* 375–377 (2004) 213–218. doi:10.1016/j.msea.2003.10.257.
- [16] J.-W. Yeh, S.-K. Chen, S.-J. Lin, J.-Y. Gan, T.-S. Chin, T.-T. Shun, C.-H. Tsau, S.-Y. Chang, Nanostructured high-entropy alloys with multiple principal elements: Novel alloy design concepts and outcomes, *Advanced Engineering Materials* 6 (5) (2004) 299–303. doi:10.1002/adem.200300567.
- [17] E. P. George, D. Raabe, R. O. Ritchie, High-entropy alloys, *Nature Reviews Materials* 4 (8) (2019) 515–534. doi:10.1038/s41578-019-0121-4.
- [18] D. B. Miracle, O. N. Senkov, A critical review of high entropy alloys and related concepts, *Acta Materialia* 122 (2017) 448–511. doi:10.1016/j.actamat.2016.08.081.
- [19] Z. Rao, P.-Y. Tung, R. Xie, Y. Wei, H. Zhang, A. Ferrari, T. P. C. Klaver, F. Körmann, P. T. Sukumar, A. Kwiatkowski da Silva, Y. Chen, Z. Li, D. Ponge, J. Neugebauer, O. Gutfleisch, S. Bauer, D. Raabe, Machine learning-enabled high-entropy alloy discovery, *Science* 378 (6615) (2022) 78–85. doi:10.1126/science.abo4940.
- [20] T. Hastings, M. Mulukutla, D. Khatamsaz, D. Salas, W. Xu, D. Lewis, N. Person, M. Skokan, B. Miller, J. Paramore, B. Butler, D. Allaire, V. Attari, I. Karaman, G. Pharr, A. Srivastava, R. Arróyave, Accelerated multi-objective alloy discovery through efficient Bayesian methods: Application to the FCC high entropy alloy space, *Acta Materialia* 297 (2025) 121173. doi:10.1016/j.actamat.2025.121173.
- [21] M. Mulukutla, A. N. Person, S. Voigt, L. Kuettner, B. Kappes, D. Khatamsaz, R. Robinson, D. Salas Mula, W. Xu, D. Lewis, H. Eoh, K. Xiao, H. Wang, J. S. Saini, R. Mahat, T. Hastings, M. Skokan, V. Attari, M. Elverud, J. D. Paramore, B. Butler, K. Vecchio, S. R. Kalidindi, D. Allaire, I. Karaman, E. L. Thomas, G. Pharr, A. Srivastava, R. Arróyave, Illustrating an effective workflow for accelerated materials discovery, *Integrating Materials and Manufacturing Innovation* 13 (2024) 453–473. doi:10.1007/s40192-024-00357-3.
- [22] M. Mulukutla, D. Khatamsaz, T. Hastings, W. Xu, D. Salas, J. Kundu, V. C. Shunmugasamy, D. Lewis, J. Hempel, C. Strosser, A. Salinas, B. Vela, N. Flores, S. Hossein Zadeh, A. Rachidi, D. Elbert, B. Butler, J. Paramore, D. Lagoudas, J. Wilkerson, G. Pharr, D. Allaire, V. Attari, I. Karaman, A. Srivastava, R. Arróyave, A BRAVE alloy design campaign (Bayesian risk-aware alloy discovery and exploration) (2026). arXiv:2606.27522, doi:10.48550/arXiv.2606.27522. URL <https://arxiv.org/abs/2606.27522>
- [23] S. P. Padhy, V. Chaudhary, Y.-F. Lim, R. Zhu, M. Thway, K. Hippalgaonkar, R. V. Ramanujan, Experimentally validated inverse design of multi-property fe-co-ni alloys, *Isience* 27 (5) (2024).
- [24] S. P. Padhy, K. P. Davidson, L. P. Tan, V. B. Varma, V. K. Sharma, X. Tan, Y. Wei, X. Xu, K. Hippalgaonkar, M. H.

- Jhon, et al., Integrated design framework for titanium aluminides through interpretable machine learning, *Journal of Alloys and Compounds* (2025) 184937.
- [25] B. Meredig, E. Antono, C. Church, M. Hutchinson, J. Ling, S. Paradiso, B. Blaiszik, I. Foster, B. Gibbons, J. Hattrick-Simpers, A. Mehta, L. Ward, Can machine learning identify the next high-temperature superconductor? Examining extrapolation performance for materials discovery, *Molecular Systems Design & Engineering* 3 (5) (2018) 819–825. doi:10.1039/C8ME00012C.
- [26] A. Y.-T. Wang, R. J. Murdock, S. K. Kauwe, A. O. Oliynyk, A. Gurlo, J. Bragoch, K. A. Persson, T. D. Sparks, Machine learning for materials scientists: An introductory guide toward best practices, *Chemistry of Materials* 32 (12) (2020) 4954–4965. doi:10.1021/acs.chemmater.0c01907.
- [27] Q. Li, N. Fu, S. S. Omeel, J. Hu, MD-HIT: Machine learning for material property prediction with dataset redundancy control, *npj Computational Materials* 10 (2024). doi:10.1038/s41524-024-01426-z.
- [28] C. K. H. Borg, C. Frey, J. Moh, T. M. Pollock, S. Gorsse, Expanded dataset of mechanical properties and observed phases of multi-principal element alloys, *Scientific Data* 7 (2020) 430. doi:10.1038/s41597-020-00768-9.
- [29] X. Yang, Y. Zhang, Prediction of high-entropy stabilized solid-solution in multi-component alloys, *Materials Chemistry and Physics* 132 (2–3) (2012) 233–238. doi:10.1016/j.matchemphys.2011.11.021.
- [30] S. Guo, C. Ng, J. Lu, C. T. Liu, Effect of valence electron concentration on stability of fcc or bcc phase in high entropy alloys, *Journal of Applied Physics* 109 (10) (2011) 103505. doi:10.1063/1.3587228.
- [31] H. Akaike, A new look at the statistical model identification, *IEEE Transactions on Automatic Control* 19 (6) (1974) 716–723. doi:10.1109/TAC.1974.1100705.
- [32] G. Schwarz, Estimating the dimension of a model, *The Annals of Statistics* 6 (2) (1978) 461–464. doi:10.1214/aos/1176344136.
- [33] A. Vehtari, A. Gelman, J. Gabry, Practical bayesian model evaluation using leave-one-out cross-validation and WAIC, *Statistics and Computing* 27 (5) (2017) 1413–1432. doi:10.1007/s11222-016-9696-4.
- [34] Y. Yao, A. Vehtari, D. Simpson, A. Gelman, Using stacking to average Bayesian predictive distributions (with discussion), *Bayesian Analysis* 13 (3) (2018) 917–1007. doi:10.1214/17-ba1091.
- [35] S. P. Padhy, S. Saurabh, K. Choudhary, et al., Automated regression workflow for interpretable deflection prediction in bio-inspired laminated composite plates, *Frontiers of Structural and Civil Engineering* 19 (2025) 1651–1668. doi:10.1007/s11709-025-1238-8.
- [36] S. Saurabh, S. P. Padhy, V. N. V. Hoang, R. Kiran, V. S. Chauhan, Automated and explainable machine learning for accelerating nonlinear vibration prediction in bio-inspired helicoidal laminated composite structures, *Structures* 86 (2026) 111358. doi:10.1016/j.istruc.2026.111358.
- [37] T. Akiba, S. Sano, T. Yanase, T. Ohta, M. Koyama, Optuna: A next-generation hyperparameter optimization framework, in: *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2019.
- [38] M. Cranmer, Interpretable machine learning for science with PySR and SymbolicRegression.jl, *arXiv preprint arXiv:2305.01582* (2023). doi:10.48550/arXiv.2305.01582.
- [39] R. Ouyang, S. Curtarolo, E. Ahmetcik, M. Scheffler, L. M. Ghiringhelli, SISSO: A compressed-sensing method for identifying the best low-dimensional descriptor in an immensity of offered candidates, *Physical Review Materials* 2 (8) (2018) 083802. doi:10.1103/PhysRevMaterials.2.083802.
- [40] M. Schneider, G. Laplanche, Data compilation regarding the effects of grain size and temperature on the strength of the single-phase FCC CrFeNi medium-entropy alloy, *Data in Brief* 34 (2021) 106712. doi:10.1016/j.dib.2020.106712.
- [41] Y.-C. Huang, C.-H. Su, S.-K. Wu, C. Lin, A study on the Hall–Petch relationship and grain growth kinetics in FCC-structured high/medium entropy alloys, *Entropy* 21 (3) (2019) 297. doi:10.3390/e21030297.
- [42] S. Yoshida, T. Ikeuchi, T. Bhattacharjee, Y. Bai, A. Shibata, N. Tsuji, Effect of elemental combination on friction stress and Hall-Petch relationship in face-centered cubic high/medium entropy alloys, *Acta Materialia* 171 (2019) 201–215. doi:10.1016/j.actamat.2019.04.017.
- [43] Z. Wu, Y. Gao, H. Bei, Thermal activation mechanisms and Labusch-type strengthening analysis for a family of high-entropy and equiatomic solid-solution alloys, *Acta Materialia* 120 (2016) 108–119. doi:10.1016/j.actamat.2016.08.047.
- [44] R. B. Figueiredo, M. Kawasaki, T. G. Langdon, Seventy years of Hall-Petch, ninety years of superplasticity and a generalized approach to the effect of grain size on flow stress, *Progress in Materials Science* 137 (2023) 101131. doi:10.1016/j.pmatsci.2023.101131.
- [45] Z. Li, S. Zhao, R. O. Ritchie, M. A. Meyers, Mechanical properties of high-entropy alloys with emphasis on face-centered cubic alloys, *Progress in Materials Science* 102 (2019) 296–345. doi:10.1016/j.pmatsci.2018.12.003.